

Midterm Project — Due Thur, 1st April

- *Submit your homework on Canvas.
- *No late homework will be accepted for credit.
- *Append the codes you used to your submission.

Problem 1: Speed and Stopping Distances of Cars

Dataset: Cars.csv

Description: The data give the speed of cars and the distances taken to stop. Note that the data were recorded in the 1920s. The dataset has 50 observations, each observation with 2 variables.

speed: numerical value describing the speed of car

dist: numerical value describing the stopping distance

Goal: Use Cross-Validation to find the relation between stopping distance and speed, assuming that *dist* is a polynomial function of speed.

1. Use leave-one-out cross validation, and plot the CV errors *versus* degree of polynomial. Report your finding and conclusion.
2. Use 5-fold cross validation, and plot the CV errors *versus* degree of polynomial. Report your finding and conclusion.
3. Compare the above two approaches and report your findings.

Problem 2: Titanic – Survival or Not

Dataset: titanic.csv

Description: The dataset contains information of passengers on the famous Titanic. It has 891 observations and each observation has 12 variables. The meaning of these variables are as follows.

survival: whether this passenger died or survived, 0=No, 1=Yes.

pclass: ticket class, 1=1st, 2=2nd, 3=3rd

sex: male or female

Age: age in years

sibsp: # of siblings/spouses aboard the Titanic

parch: # of parents/children aboard the Titanic

ticket: Ticket number

fare: Passenger fare

cabin: Cabin number

embarked: Port of Embarkation, C=Cherbourg, Q=Queenstown, S=Southampton

Goal: apply Bootstrap for the statistical inference of how these variables affect the probability of survival.

1. Treat *survival* as response, fit a logistic regression model using predictors *pclass*, *sex*, *age*, *sibsp* and *fare*. Report the estimated coefficients for *Sex/male* and *pclass/3rd*, and report their 95% confidence intervals.

2. Now, apply Bootstrap with 1000 repetitions to obtain the 95% confidence intervals for the above coefficients. How do they compare with the reported confidence intervals from above.
3. Explore the dataset as you like and report some of your findings.

Problem 3: Predicting First-Year College Students' GPA

Dataset: FirstYearGPA.csv

Description: The dataset contains information of 219 first-year college students. Each student has 10 variables including his/her GPA in the first year of college. The meaning of these variables are as follows.

GPA: First-year college GPA on a 0.0 to 4.0 scale.

HSGPA: High school GPA on a 0.0 to 4.0 scale

SATV: Verbal/critical reading SAT score

SATM: Math SAT score

Male: 1=male, 0=female

HU: Number of credit hours earned in humanities courses in high school

SS: Number of credit hours earned in social science courses in high school

FirstGen: 1= student is the first in her or his family to attend college; 0=otherwise

White: 1=white students; 0=others

CollegeBound: 1=attended a high school where $\geq 50\%$ students intended to go on to college; 0=otherwise

Goal: select the best linear model to predict first-year college students' GPA.

1. Use all variables to predict students' first-year GPA, by best subset selection up to the size of 8. Report the 8 best linear models (best model for each model size $k = 1, \dots, 8$) and plot the R-square *versus* model size. Which is the best model using adjusted R-square?
2. Repeat Step 1 but choose the best model by 5-fold CV. Report the summary of that model.
3. Use all variables to predict students' first-year GPA, by forward stepwise selection up to the size of 8. Report the 8 (best model for each model size $k = 1, \dots, 8$) best linear models and plot the adjusted R-square *versus* model size. Which is the best model using BIC?
4. Repeat Step 3 but choose the best model by 5-fold CV. Report the summary of that model.

Problem 4: Prediction of the Progression of Diabetes

Dataset: diabetes_train.csv and diabetes_test.csv

Description: The dataset contains information of 442 patients of diabetes. For each patient, the dataset provides some measure of 10 baseline variables. The meaning of these variables are as follows.

Y: a quantitative measure of the progression of diabetes.

age: age of the patient (note that the column is standardized)

sex: male or female (note that they are originally coded by number 1 and 2, then the column is standardized)

bmi: body mass index (column is standardized)

map: some numerical measure of blood

tc: some numerical measure of blood

ldl: some numerical measure of blood

hdl: some numerical measure of blood

tch: some numerical measure of blood

ltg: some numerical measure of blood

glu: some numerical measure of blood

\vdots

age.tch: the interaction of baseline variables *age* and *tch*

\vdots

Goal: apply regularized regression to find which baseline variables or their interactions affect the progression of diabetes.

1. Use the train dataset to fit LASSO estimators with regularization parameter λ chosen from the grid $10^{seq(4,-2,length=100)}$. Plot the coefficients *versus* the ℓ_1 norm. Report any findings you think interesting.
2. Use train dataset to fit LASSO and apply 10-fold cross validation. Plot the CV error *versus* the values of lambda. What is the best lambda value according CV error? Report the linear model using the best lambda value, how many variables are included in the model, what are they?
3. Use the best model in Step 2 to predict the progression of diabetes on the test dataset, and report the mean test error.