


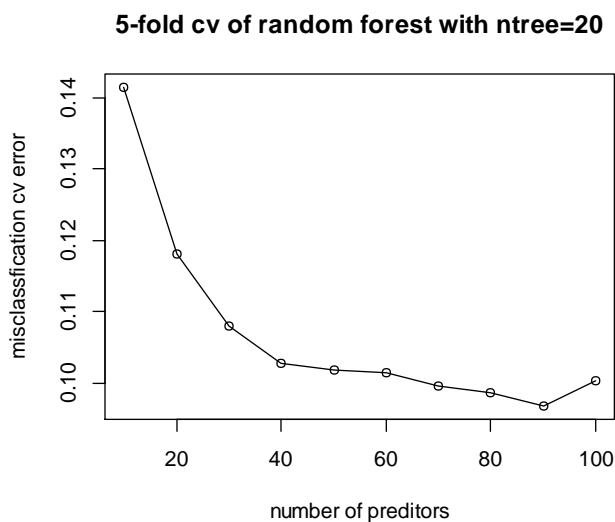
Part I. Classification on 20newsgroup Data

1. Fixing $mtry = \sqrt{100} = 10$ with different $ntree$ (2 to 50, step = 2), the result is:



We can see that the misclassification cv error decreases as number of tree increase until $ntree = 20$, and the error does not vary much for $ntree > 20$.

Fixing $ntree = 20$ with different $mtry$ (10 to 100, step=10), the result is:



The best CV error was 0.0967375 with $ntree = 20$, $mtry = 90$.

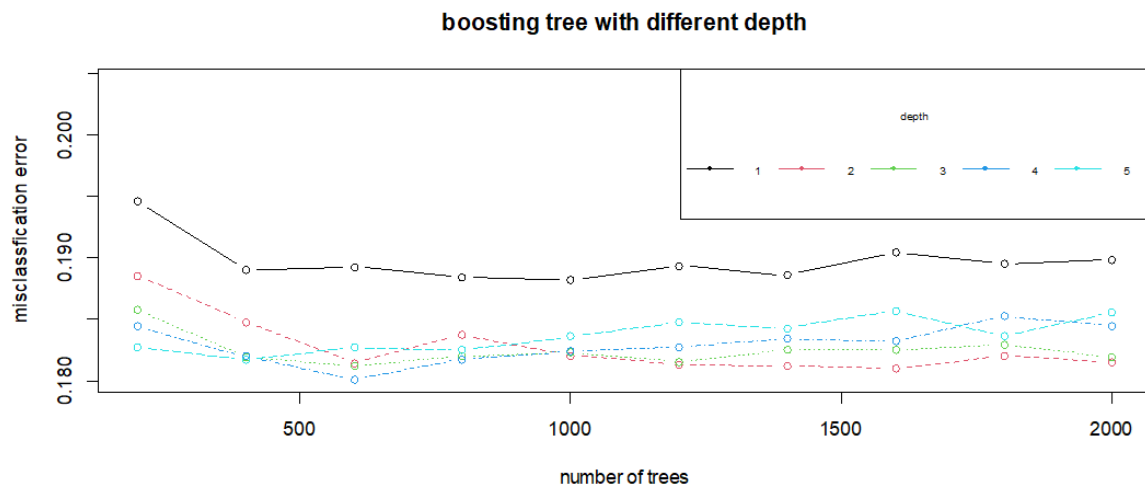
A random forest model is retrained with $ntree=20$ and $mtry=90$, the misclassification error is 0.09826376 and the corresponding confusion matrix is:

		newsgroup			
predicted_newsgroup		1	2	3	4
1	4367	205	271	158	
2	90	3096	119	131	
3	73	78	2111	100	
4	75	140	156	5072	

The higher the Gini index, the more important the variable is. The 10 most important keywords are:

Keywords	Gini Index
windows	720.1825
god	530.8854
government	443.3101
team	418.4419
car	373.9237
christian	371.5577
jews	262.3375
space	239.3162
baseball	234.6961
graphics	216.2313

2. Boosting trees were built with `ntree=2000`, `interaction.depth` from 1 to 5 and the default `shrinkage=0.1` using 5-fold cross validation. The misclassification cv error are plotted as below:



Best CV error was 0.18 with `ntree=600` and `interaction.depth=4`.

A boosting tree model is retrained with `ntree=600`, `interaction.depth=4` and `shrinkage=0.1`, the misclassification error is 0.1520749 and the corresponding confusion matrix is:

		newsgroup			
		1	2	3	4
predicted_newsgroup	1	4220	254	476	216
	2	63	2821	96	101
	3	153	129	1781	194
	4	169	315	304	4950

3. In general, boosting tree model outperforms random forest when it is well tuned. However, for this dataset, random forest is a better model than boosting tree. One reason is that the data contains much noise which caused boosting tree model to over fit.

4. Building a multi-class LDA classifier, the 5-fold CV error of misclassification is 0.202045. A multi-class LDA classifier is retrained with all data and the confusion matrix is:

		newsgroup			
predicted_newsgroup		1	2	3	4
1	4102	341	620	268	
2	45	2607	119	125	
3	239	181	1530	299	
4	219	390	388	4769	

5. When building a multi-class QDA classifier with all variables, a rank deficiency error is given. The error was due to the covariance matrix is not invertible in some groups. It is observed that values of certain variables in certain group are all 0, which means all records of certain group do not contain such keywords. To fix this error, we should not include any keywords that are all 0 in particular group in the qda model. The following keywords are identified from the groups of each folds of the cross validation that should be removed:

```
> unique(word_not_use)
[1] "bmw"      "hockey"    "jews"      "lunar"     "nhl"       "puck"      "vitamin"   "aids"
[9] "bible"    "dos"       "israel"    "orbit"     "patients"  "scsi"      "shuttle"   "honda"
[17] "jesus"    "mars"      "dealer"    "season"
```

Thus, building a multi-class QDA classifier with all but the above variables, the 5-fold CV error of misclassification is 0.2738134. A multi-class LDA classifier is retrained with all data and the confusion matrix is:

		newsgroup			
predicted_newsgroup		1	2	3	4
1	4246	753	1075	738	
2	133	2464	217	432	
3	149	88	1058	198	
4	77	214	307	4093	

6. Both LDA and QDA did not perform better than random forest and boosting tree. LDA has large bias while some important variables (e.g. "jews") could not be used in QDA due to its limitation which lowered the prediction accuracy.

Part II. Spectral Clustering on 20newsgroup Data

1. The mis-clustering error rate using top 4 left singular vectors from PCA and K-means with K=4 is $1 - \frac{4567+1056+336+2102}{16242} = 0.503694$.
2. The mis-clustering error rate using top 5 left singular vectors from PCA and K-means with K=4 is $1 - \frac{4573+1054+338+2247}{16242} = 0.494397$.

```
> doc.pc=prcomp(doc.table, scale=TRUE)
> set.seed(1)
> doc.km4=kmeans(doc.pc$x[,1:4],4,nstart=20)
> require(plyr)
> doc.km4$cluster=mapvalues(doc.km4$cluster,from=c(1,2,3,4),to=c(2,1,4,3))
> table(doc.km4$cluster,doc.group[,1])
```

	1	2	3	4
1	4567	2418	1996	3341
2	3	1056	5	7
3	18	0	336	11
4	17	45	320	2102

```
> doc.km5=kmeans(doc.pc$x[,1:5],4,nstart=20)
> doc.km5$cluster=mapvalues(doc.km5$cluster,from=c(1,2,3,4),to=c(4,3,2,1))
> table(doc.km5$cluster,doc.group[,1])
```

	1	2	3	4
1	4573	2429	2138	3195
2	3	1054	5	8
3	18	0	338	11
4	11	36	176	2247

3. The performances using PCA with the top 4 and 5 left singular vectors for clustering is much worse than the method used from part I. One reason is that the top 4 and top 5 only account for 13.7% and 16% of total variance of the data. Another reason is simply K-means is not suitable for clustering in this dataset which may be due to the way that K-means measuring the distance does not align with how the groups are separated.

```
> summary(doc.pc)
Importance of components:
```

	PC1	PC2	PC3	PC4	PC5
Standard deviation	2.00581	1.92846	1.76839	1.68292	1.53043
Proportion of Variance	0.04023	0.03719	0.03127	0.02832	0.02342
Cumulative Proportion	0.04023	0.07742	0.10869	0.13702	0.16044

Part III. Classification on MNIST Data

1. Using 5-fold cross validation of linear kernel, the cost parameter has been tuned for 53.7 seconds in which the best cost is 0.1.

```
> mnist.train=read.csv("train_resized.csv",header=TRUE)
> mnist.test=read.csv("test_resized.csv",header=TRUE)
> mnist.train$label=as.factor(mnist.train$label)
> mnist.test$label=as.factor(mnist.test$label)
> mnist.train.i36=mnist.train[mnist.train[,1]==3 | mnist.train[,1]==6,]
> mnist.test.i36=mnist.test[mnist.test[,1]==3 | mnist.test[,1]==6,]
> library(e1071)
> set.seed(1)
> tc=tune.control(cross = 5)
> start_time=Sys.time()
> tune.out=tune(svm,label~.,data=mnist.train.i36,kernel="linear",tunecontrol=tc,ranges=list(cost=c(0.001, 0.01, 0.1, 1,5,10,100)))
> end_time=Sys.time()
> end_time-start_time
Time difference of 53.71672 secs
> summary(tune.out)
```

Parameter tuning of 'svm':

- sampling method: 5-fold cross validation

- best parameters:

cost
0.01

- best performance: 0.003816051

- Detailed performance results:

	cost	error	dispersion
1	1e-03	0.005641777	0.002712912
2	1e-02	0.003816051	0.002594954
3	1e-01	0.005475665	0.003190432
4	1e+00	0.007964878	0.002660070
5	5e+00	0.008130853	0.002444353
6	1e+01	0.008130853	0.002444353
7	1e+02	0.008130853	0.002444353

The misclassification error on the test data is $\frac{6+11}{1251+6+11+1194} = 0.69\%$ and the confusion matrix is:

```
> bestmod=tune.out$best.model
> ypred=predict(bestmod ,mnist.test.i36)
> table(predict=ypred, truth=mnist.test.i36$label)
```

[illegible]

- Using 5-fold cross validation of radial kernel, the cost and gamma parameters have been tuned for 12 mins in which the best (cost, gamma) pair is (100, 0.0001).

```
> start_time=Sys.time()
> tune.out_1=tune(svm, label~. ,data=mnist.train.i36,kernel="radial",tunecontrol=tc,ranges=list(cost=c(1,10,100,1000),gamma=c(0.0001,0.001,0.01,0.1)))
> end_time=Sys.time()
> end_time-start_time
Time difference of 12.02154 mins
> summary(tune.out_1)
```

Parameter tuning of 'svm':

- sampling method: 5-fold cross validation

- best parameters:

```
cost gamma
100 1e-04
```

- best performance: 0.004314802

- Detailed performance results:

	cost	gamma	error	dispersion
1	1	1e-04	0.008961417	0.001486086
2	1	1e-04	0.005144540	0.001799600
3	100	1e-04	0.004314802	0.002517410
4	1000	1e-04	0.006305815	0.002390097
5	1	1e-03	0.005310790	0.002391457
6	10	1e-03	0.004646890	0.002346419
7	100	1e-03	0.005144815	0.003776069
8	1000	1e-03	0.004978978	0.003936735
9	1	1e-02	0.009127943	0.004066597
10	10	1e-02	0.008463905	0.002899872
11	100	1e-02	0.008463905	0.002899872
12	1000	1e-02	0.008463905	0.002899872
13	1	1e-01	0.123591675	0.002932574
14	10	1e-01	0.123631221	0.001344517
15	100	1e-01	0.123631221	0.001344517
16	1000	1e-01	0.123631221	0.001344517

The misclassification error on the test data is $\frac{4+10}{1252+4+10+1196} = 0.57\%$ and the confusion matrix is:

```
> bestmod_r1=tune.out_1$best.model
> ypred=predict(bestmod_r1 ,mnist.test.i36)
> table(predict=ypred, truth=mnist.test.i36$label)
```

[illegible]

3. Both models above perform equally well for this problem when the parameters are carefully tuned while it takes longer to tune the model with radial kernel.
4. Using 5-fold cross validation of linear kernel, the cost parameter has been tuned for 9 mins in which the best cost is 0.1.

```
> mnist.train.i1258=mnist.train[mnist.train[,1]==1 | mnist.train[,1]==2 | mnist.train[,1]==5 | mnist.train[,1]==8,]
> mnist.test.i1258=mnist.test[mnist.test[,1]==1 | mnist.test[,1]==2 | mnist.test[,1]==5 | mnist.test[,1]==8,]
> start_time=Sys.time()
> tune.out.1258=tune(svm,label~.,data=mnist.train.i1258,kernel="linear",tunecontrol=tc,ranges=list(cost=c(0.01,0.1,1,10,100)))
> end_time=Sys.time()
> end_time-start_time
Time difference of 8.997765 mins
> summary(tune.out.1258)
```

Parameter tuning of 'svm':

```
- sampling method: 5-fold cross validation

- best parameters:
  cost
  0.1

- best performance: 0.03970427

- Detailed performance results:
  cost      error dispersion
1 1e-02 0.04180286 0.004959515
2 1e-01 0.03970427 0.004464172
3 1e+00 0.04348173 0.003614521
4 1e+01 0.04902195 0.005449180
5 1e+02 0.05179178 0.004165644
```

The misclassification error on the test data is $1 - \frac{1344+1140+1062+1041}{4806} = 4.56\%$ and the confusion matrix is:

```
> bestmod_1258=tune.out.1258$best.model
> ypred=predict(bestmod_1258 ,mnist.test.i1258)
> table(predict=ypred, truth=mnist.test.i1258$label)
```

	truth									
predict	0	1	2	3	4	5	6	7	8	9
0	0	0	0	0	0	0	0	0	0	0
1	0	1344	5	0	0	12	0	0	19	0
2	0	11	1140	0	0	19	0	0	26	0
3	0	0	0	0	0	0	0	0	0	0
4	0	0	0	0	0	0	0	0	0	0
5	0	2	16	0	0	1062	0	0	46	0
6	0	0	0	0	0	0	0	0	0	0
7	0	0	0	0	0	0	0	0	0	0
8	0	6	24	0	0	33	0	0	1041	0
9	0	0	0	0	0	0	0	0	0	0

5. Using 5-fold cross validation of linear kernel on full dataset, the cost parameter has been tuned for 56 mins in which the best cost is 0.1.

```

> start_time=Sys.time()
> options(warn=-1)
> tune.out.all=tune(svm,label=.,data=mnist.train,kernel="linear",tunecontrol=tc,ran
ges=list(cost=c(0.01,0.1,1,10,100)))
> end_time=Sys.time()
> end_time-start_time
Time difference of 55.55482 mins
> summary(tune.out.all)

```

Parameter tuning of 'svm':

- sampling method: 5-fold cross validation

- best parameters:

```
cost
0.1
```

- best performance: 0.0626

- Detailed performance results:

```

cost      error dispersion
1 1e-02 0.06373333 0.003079051
2 1e-01 0.06260000 0.003001389
3 1e+00 0.06903333 0.003795831
4 1e+01 0.07366667 0.003717451
5 1e+02 0.07656667 0.003724170

```

The misclassification error on the test data is

$$1 - \frac{1108+1341+1112+1136+1115+1023+1148+1183+1004+1058}{12000} = 6.43\% \text{ and the confusion}$$

matrix is:

```

> bestmod_all=tune.out.all$best.model
> ypred=predict(bestmod_all ,mnist.test)
> table(predict=ypred, truth=mnist.test$label)

```

	truth	0	1	2	3	4	5	6	7	8	9
predict	0	1108	0	5	1	2	7	10	1	5	4
	1	0	1341	2	7	3	12	3	3	16	6
	2	2	8	1112	38	12	6	12	13	22	8
	3	1	2	5	1136	0	33	0	4	24	9
	4	1	0	24	0	1115	3	8	7	5	26
	5	10	2	3	39	3	1023	13	5	35	6
	6	10	1	6	3	7	17	1148	0	2	0
	7	0	2	9	7	5	1	1	1183	3	29
	8	6	6	16	19	1	18	4	2	1004	7
	9	2	1	3	12	27	6	1	46	16	1058

Part IV. Additional Bonus

Neural network could be used for this dataset. A Multi-Layer Perceptron (MLP) with 6 hidden layer was trained as below:

```

> library(dplyr)
> library(keras)
> library(tensorflow)
> library(yardstick)
>
> data_train=read.csv("train_resized.csv", header=TRUE)
> data_test=read.csv("test_resized.csv", header=TRUE)
>
> train_x=as.matrix(data_train[,-1])
> test_x=as.matrix(data_test[,-1])
> train_y=to_categorical(data_train$label, num_classes = 10)
> test_y=data_test$label
>
> tf$random$set_seed(1)
> start_time=Sys.time()
> model=keras_model_sequential(name = "MLP_MNIST")
> layer_dense(model,units = 512, activation = "relu", input_shape = ncol(train_x), name = "Hidden_1")
> layer_dense(model,units = 256, activation = "relu", name = "Hidden_2")
> layer_dense(model,units = 128, activation = "relu", name = "Hidden_3")
> layer_dense(model,units = 64, activation = "relu", name = "Hidden_4")
> layer_dense(model,units = 32, activation = "relu", name = "Hidden_5")
> layer_dense(model,units = 16, activation = "relu", name = "Hidden_6")
> layer_dense(model,units = 10, activation = "softmax", name = "out")

```

Model

Model: "MLP_MNIST"

Layer (type)	Output shape	Param #
Hidden_1 (Dense)	(None, 512)	74240
Hidden_2 (Dense)	(None, 256)	131328
Hidden_3 (Dense)	(None, 128)	32896
Hidden_4 (Dense)	(None, 64)	8256
Hidden_5 (Dense)	(None, 32)	2080
Hidden_6 (Dense)	(None, 16)	528
out (Dense)	(None, 10)	170

Total params: 249,498


Trainable params: 249,498

Non-trainable params: 0

```

> compile(model, loss="categorical_crossentropy", optimizer=optimizer_adam(lr=0.001), metrics="accuracy")
> history=fit(model, x=train_x, y=train_y, epochs=15, batch_size=32, validation_split=0.2, verbose=1)
Epoch 1/15
750/750 [=====] - 3s 3ms/step - loss: 1.1758 - accuracy: 0.7621 - val_loss: 0.2616 - val_accuracy: 0.9278
Epoch 2/15
750/750 [=====] - 3s 3ms/step - loss: 0.1966 - accuracy: 0.9442 - val_loss: 0.1845 - val_accuracy: 0.9467
Epoch 3/15
750/750 [=====] - 2s 3ms/step - loss: 0.1408 - accuracy: 0.9595 - val_loss: 0.1795 - val_accuracy: 0.9508
Epoch 4/15
750/750 [=====] - 3s 3ms/step - loss: 0.1316 - accuracy: 0.9635 - val_loss: 0.1651 - val_accuracy: 0.9527
Epoch 5/15
750/750 [=====] - 2s 3ms/step - loss: 0.1015 - accuracy: 0.9713 - val_loss: 0.1687 - val_accuracy: 0.9528
Epoch 6/15
750/750 [=====] - 3s 3ms/step - loss: 0.0836 - accuracy: 0.9753 - val_loss: 0.1652 - val_accuracy: 0.9555
Epoch 7/15
750/750 [=====] - 2s 3ms/step - loss: 0.0828 - accuracy: 0.9762 - val_loss: 0.1496 - val_accuracy: 0.9633
Epoch 8/15
750/750 [=====] - 3s 3ms/step - loss: 0.0679 - accuracy: 0.9807 - val_loss: 0.1533 - val_accuracy: 0.9643
Epoch 9/15
750/750 [=====] - 2s 3ms/step - loss: 0.0653 - accuracy: 0.9811 - val_loss: 0.1227 - val_accuracy: 0.9697
Epoch 10/15
750/750 [=====] - 3s 3ms/step - loss: 0.0534 - accuracy: 0.9852 - val_loss: 0.1334 - val_accuracy: 0.9683
Epoch 11/15
750/750 [=====] - 2s 3ms/step - loss: 0.0625 - accuracy: 0.9814 - val_loss: 0.1299 - val_accuracy: 0.9675
Epoch 12/15
750/750 [=====] - 2s 3ms/step - loss: 0.0480 - accuracy: 0.9860 - val_loss: 0.1738 - val_accuracy: 0.9613
Epoch 13/15
750/750 [=====] - 2s 3ms/step - loss: 0.0524 - accuracy: 0.9856 - val_loss: 0.1377 - val_accuracy: 0.9700
Epoch 14/15
750/750 [=====] - 2s 3ms/step - loss: 0.0413 - accuracy: 0.9881 - val_loss: 0.1422 - val_accuracy: 0.9732
Epoch 15/15
750/750 [=====] - 2s 3ms/step - loss: 0.0520 - accuracy: 0.9867 - val_loss: 0.1416 - val_accuracy: 0.9685
> end_time=Sys.time()
> end_time-start_time
Time difference of 39.05762 secs
> plot(history)

```

The misclassification error of the test data is 2.9% and the training time of the model is 39 seconds. The confusion matrix is:

```
> pred_test=predict_classes(model, test_x)
> error_mlp=1-accuracy_vec(truth = as.factor(data_test$label), estimate = as.factor(pred_test))
> error_mlp
[1] 0.029
> table(prediction = as.factor(pred_test), truth = as.factor(data_test$label))
```


	truth	0	1	2	3	4	5	6	7	8	9
prediction	0	1120	0	4	0	3	3	2	1	1	3
	1	0	1340	2	1	1	2	4	1	11	2
	2	0	9	1148	4	0	0	1	6	5	2
	3	1	3	7	1234	0	25	0	0	20	5
	4	0	2	3	0	1125	0	2	1	1	8
	5	2	0	0	6	1	1075	9	0	8	2
	6	7	0	1	0	7	4	1177	0	1	0
	7	0	4	13	10	4	1	0	1242	5	13
	8	7	5	6	3	0	9	3	0	1075	2
	9	3	0	1	4	34	7	2	13	5	1116

The performance of neural network outperforms svm in part III. However, there are much more hyper-parameters need to be set for neural network (like number of hidden layers, number of neurons in each layer, activation function etc.) and it is hard to reach the optimal set of hyper-parameters. I have also trained a CNN model but the performance is similar to that of the MLP. The reason may be the dataset is not large enough.

```
> train_x_cnn=array_reshape(train_x, dim=c(nrow(train_x), 12, 12, 1))
> test_x_cnn=array_reshape(test_x, dim=c(nrow(test_x), 12, 12, 1))
>
> tf$random$set_seed(1)
> model1=keras_model_sequential(name = "CNN_Mnist")
> layer_conv_2d(model1,filters=32, kernel_size=c(2,2), padding="same", activation="relu", input_shape=c(12,12,1))
> layer_max_pooling_2d(model1, pool_size=c(2,2))
> layer_conv_2d(model1, filters=32, kernel_size=c(2,2), padding="same", activation="relu", input_shape=c(12,12,1))
> layer_max_pooling_2d(model1, pool_size=c(2,2))
> layer_conv_2d(model1, filters=32, kernel_size=c(2,2), padding="same", activation="relu", input_shape=c(12,12,1))
> layer_max_pooling_2d(model1, pool_size=c(2,2))
> layer_flatten(model1)
> layer_dense(model1,units=16, activation="relu")
> layer_dense(model1,units=10, activation="softmax", name="Output")
> model1
Model
Model: "CNN_Mnist"
```

Layer (type)	Output Shape	Param #
conv2d_3 (Conv2D)	(None, 12, 12, 32)	160
max_pooling2d_3 (MaxPooling2D)	(None, 6, 6, 32)	0
conv2d_4 (Conv2D)	(None, 6, 6, 32)	4128
max_pooling2d_4 (MaxPooling2D)	(None, 3, 3, 32)	0
conv2d_5 (Conv2D)	(None, 3, 3, 32)	4128
max_pooling2d_5 (MaxPooling2D)	(None, 1, 1, 32)	0
flatten_1 (Flatten)	(None, 32)	0
dense_1 (Dense)	(None, 16)	528
Output (Dense)	(None, 10)	170
Total params: 9,114		
Trainable params: 9,114		
Non-trainable params: 0		

```
> compile(model1, loss="categorical_crossentropy", optimizer=optimizer_adam(lr=0.001), metrics="accuracy")
> history=fit(model1, x=train_x_cnn, y=train_y, epochs=15, batch_size=32, validation_split=0.2, verbose=1)
Epoch 1/15
750/750 [=====] - 7s 8ms/step - loss: 2.1869 - accuracy: 0.3098 - val_loss: 0.6360 - val_accuracy: 0.8702
Epoch 2/15
750/750 [=====] - 6s 8ms/step - loss: 0.5400 - accuracy: 0.8891 - val_loss: 0.3081 - val_accuracy: 0.9188
Epoch 3/15
750/750 [=====] - 7s 9ms/step - loss: 0.2684 - accuracy: 0.9244 - val_loss: 0.2251 - val_accuracy: 0.9408
Epoch 4/15
750/750 [=====] - 7s 9ms/step - loss: 0.2027 - accuracy: 0.9443 - val_loss: 0.2769 - val_accuracy: 0.9177
Epoch 5/15
750/750 [=====] - 7s 9ms/step - loss: 0.1698 - accuracy: 0.9520 - val_loss: 0.1983 - val_accuracy: 0.9440
Epoch 6/15
750/750 [=====] - 9s 12ms/step - loss: 0.1423 - accuracy: 0.9596 - val_loss: 0.1752 - val_accuracy: 0.9537
Epoch 7/15
750/750 [=====] - 8s 11ms/step - loss: 0.1154 - accuracy: 0.9666 - val_loss: 0.1554 - val_accuracy: 0.9552
Epoch 8/15
750/750 [=====] - 9s 12ms/step - loss: 0.1000 - accuracy: 0.9689 - val_loss: 0.1566 - val_accuracy: 0.9547
Epoch 9/15
750/750 [=====] - 9s 12ms/step - loss: 0.0922 - accuracy: 0.9714 - val_loss: 0.1466 - val_accuracy: 0.9592
Epoch 10/15
750/750 [=====] - 7s 10ms/step - loss: 0.0806 - accuracy: 0.9736 - val_loss: 0.1328 - val_accuracy: 0.9610
Epoch 11/15
750/750 [=====] - 9s 12ms/step - loss: 0.0772 - accuracy: 0.9756 - val_loss: 0.1395 - val_accuracy: 0.9585
Epoch 12/15
750/750 [=====] - 6s 9ms/step - loss: 0.0739 - accuracy: 0.9786 - val_loss: 0.1376 - val_accuracy: 0.9612
Epoch 13/15
750/750 [=====] - 9s 12ms/step - loss: 0.0716 - accuracy: 0.9781 - val_loss: 0.1136 - val_accuracy: 0.9692
Epoch 14/15
750/750 [=====] - 9s 13ms/step - loss: 0.0640 - accuracy: 0.9804 - val_loss: 0.1321 - val_accuracy: 0.9647
Epoch 15/15
750/750 [=====] - 7s 9ms/step - loss: 0.0572 - accuracy: 0.9827 - val_loss: 0.1193 - val_accuracy: 0.9673
> plot(history)
```



```
> pred_test=predict_classes(model1, test_x_cnn)
> error_cnn=1-accuracy_vec(truth = as.factor(data_test$label), estimate = as.factor(pred_test))
> error_cnn
[1] 0.03141667
> table(prediction = as.factor(pred_test), truth = as.factor(data_test$label))
```

	truth	0	1	2	3	4	5	6	7	8	9
prediction	0	1119	0	7	0	4	1	1	1	8	5
	1	0	1347	3	2	3	0	1	2	9	1
	2	3	6	1145	21	3	0	1	5	8	1
	3	0	2	2	1198	0	13	0	5	4	1
	4	0	1	3	0	1134	0	4	4	1	11
	5	0	0	0	7	0	1079	1	0	6	6
	6	13	2	0	0	3	19	1188	0	3	0
	7	0	0	8	8	0	0	0	1211	0	3
	8	5	5	16	21	0	13	4	4	1087	10
	9	0	0	1	5	28	1	0	32	6	1115